

CSE231 – Operating Systems

Assignment 3

Writeup

Name: Ansh Arora

Roll No: 2019022

Task: Modifying a CFS Scheduler by adding soft real-time requirement functionalities to it which have a higher priority than vruntime, and then add a system call to modify process's real time requirements.

Description of Code and Implementation –

In the question, we were asked to add a soft real time requirement to the process's task struct's schedule entity (`rt_val`). This guarantees the task at least x units of time slice. Additionally, any process that has a real time `rt_val` value given to it, that process will have priority over the process which doesn't have said `rt_val` guarantees.

Here are the changes that are made to the kernel –

1) The `rt_val` identifier is added to the sched entity of task struct of a process. This is done in the `sched.h` file where `rt_val` is added in the `u64` format.

2) The value of `rt_val` is initialized to zero in the function `__sched_fork` of `core.c` file, where other data members of the schedule entity are also initialized.

3) In the `fair.c` file of the CFS Scheduler, we have made changes to the `entity_before()` function, which

is basically a comparator that returns value 1 or 0 on basis of whether the a process needs to come on left on the RB tree or the b process, respectively. While this is traditionally done using the vruntime values, here we add functionality to check the soft real time requirements as well. For this, we add conditional statements. If both the processes a and b are having a soft time requirement, then the process with the lower soft realtime requirement is given the higher priority. If either one of the two processes is having a real time requirement, the process having the realtime requirement is given priority and is made to appear on the left side on the RB Tree. In case both a and b don't have soft guarantees, then the normal scheduling placement using vruntime is followed.

4) Another change in fair.c file is made in the function `update_curr()`. Here, we subtract the execution time of the process from the x slices of guaranteed time (`rt_val`, our soft realtime requirement) after it runs. The `delta_exec` value is decremented from the `rt_val` value if and only if the process has a `rt_val` value, for which a check has been added in the `update_curr()` function. If that isn't the case, then the normal functioning occurs where the vruntime value is decremented as it is.

The `rtnice` system call is added using the `SYSCALL_DEFINE2` function, which takes in the process in which soft realtime guarantee is to be added and the soft real time guarantee value as its parameters. The syscall uses the `find_get_pid` and `pid_task` methods to get the `task_struct` of the given pid process, and then allots it the passed real time values. Errors are handled. The successful processing of this system call is also accompanied with a kernel

message, which can be accessed from the terminal using the dmesg command.

The syscall is implemented by adding the syscall in the syscall table and also adding the folder of its presence in the core-y of the main Makefile.

Inputs Taken from User –

The syscall takes in the pid, and soft realtime values in seconds. The value is then multiplied by 10^9 inside the system call by itself, to convert it into nanoseconds. The user can run the two testfiles by the make commands as given in the below Makefile.

```
compile:
|   gcc -lgomp test.c

run: compile
|   ./a.out

compile2:
|   gcc -lgomp test2.c

run2: compile2
|   ./a.out

clear:
|   rm a.out
```

Output Expected -

In the first test file, I have shown the difference between running with and without soft real time guarantees. I create 10 processes for each, individually and make them run concurrently.

```
ansharora@ubuntu:~$ make test1
gcc test1.c -lgomp
./a.out
Without Soft Real-Time Values-
Process finished in 4.702967 seconds
Process finished in 4.935455 seconds
Process finished in 5.012862 seconds
Process finished in 5.027732 seconds
Process finished in 5.075862 seconds
Process finished in 5.075538 seconds
Process finished in 5.092636 seconds
Process finished in 5.110232 seconds
Process finished in 5.120620 seconds
Process finished in 5.134664 seconds
With Soft Real-Time Values-
Process with RT finished in 1.526738 seconds
Process with RT finished in 2.047495 seconds
Process with RT finished in 2.124027 seconds
Process with RT finished in 2.538282 seconds
Process with RT finished in 3.111094 seconds
Process with RT finished in 3.663553 seconds
Process with RT finished in 4.628839 seconds
Process with RT finished in 4.694925 seconds
Process with RT finished in 5.199664 seconds
Process with RT finished in 5.543255 seconds
```

As we can see in the attached screenshot, in the case without real time values, the time taken in seconds comes out to be much higher than it does so in the case of soft real time values, where the processes having the real time values are given the highest priority, hence being able to run much faster than in the case when no real time guarantees are passed.

In the second test file, I have run 10 processes, few of them having real time guarantees and the others not, concurrently. The process for which real time guarantees are offered are picked randomly using `rand()`. All processes run the same function once again.

```
ansharora@ubuntu:~$ make test2
gcc test2.c -lgomp
./a.out
With Mix of Non Soft-Time and Soft Real-Time Values-
7. with RT 1.533206 seconds
2. with RT 1.548328 seconds
8. with RT 1.987190 seconds
3. 4.736953 seconds
4. 4.797254 seconds
9. 4.807710 seconds
5. 4.972708 seconds
1. 4.974130 seconds
6. 5.094361 seconds
10. 5.135022 seconds_
```

In this case, as we can see, the randomly selected processes given real time guarantees (depicted using with RT) end up running much faster performance than processes that don't have these soft realtime guarantees. Thus the processes having soft realtime values are given priority over other processes.

All above processes are created as child processes inside the testfiles and exit() on termination. They all run the same function func() in the given time.

Error Values and How to Interpret Them-

1. ESRCH (Errno = 3) No such Process Found. This is generated when the given pid doesn't have a process attached to it. This is handled in the test.c file as an errno generated, and the error number is printed out along with error message using perror() and stderr().

2. EINVAL (Errno = 22) Invalid argument. This is generated when the given pid or soft real time value (syscall's arguments) lie in a range that is unacceptable to the system (when they are negative). These are again handled using perror() and stderr().

Additionally, in my test file I have created a function to check the correct functioning of the errors using error check.

```
ERROR CHECK (Y=1 or 2/N=3) 1  
Value of errno: 22  
ERROR: Invalid argument  
Error encountered: Invalid argument
```

```
ERROR CHECK (Y=1 or 2/N=3) 2  
Value of errno: 3  
ERROR: No such process  
Error encountered: No such process
```